

Games

Non-cooperative multi-agent systems

Games

Many problems can be modelled as games: multiple agents with (possibly competing) interests:

- Chess
- Go
- Multiple agents competing for limited resources

Games and agents acting in them can be described by:

- **Actions** available to each agent in various stages of a game
- **Utility** (payoff) functions, one for each agent. They assign a real number to every possible outcome of the game (typically terminal states). [If the utility functions are the same the agents become cooperative.]
- **Strategy** functions, one for each agent. They determine what action must be taken in each state. [Typically the goal is to find a strategy that maximises the utility for one agent or for a group of agents.]

Example: rock paper scissors

		Bob		
		<i>rock</i>	<i>paper</i>	<i>scissors</i>
Alice	<i>rock</i>	0,0	-1,1	1,-1
	<i>paper</i>	1,-1	0,0	-1,1
	<i>scissors</i>	-1,1	1,-1	0,0

- The game is not turn-based; it is a *simultaneous action game*.
- The game has 9 possible outcomes.
- The table shows two utility functions (Alice's and Bob's).

Game Trees

In *perfect-information games* (where agents can see which states they are in) a **game tree** is a finite tree where the nodes are states and the arcs correspond to actions by the agents. In particular:


- Each internal node is labeled with an agent (or with nature). The agent is said to control the node.
- Each arc out of a node labeled with agent i corresponds to an action for agent i .
- Each internal node labeled with **nature** has a probability distribution over its children.
- The leaves represent final outcomes and are labeled with a utility for each agent.

Perfect-information zero-sum turn-based games

Properties:

- Typically two agents
- They take turn to play
- There is no chance involved.
- The state of the game is fully observable by all agents
- Utility (payoff) values for each agent are the opposite of those of the other. Also called **zero sum**; the sum of the reward and penalty is constant.
- Because they are opposite of each other, we use one value and assume that one player tries to maximise the utility, the other minimise.
- This creates adversary.

Example: noughts and crosses




- A game tree can be seen as a contingency tree containing all possibilities
- The root is the current state
- An actual play is a path from the root to one of the leaf nodes

An optimal strategy: Min-Max function

Designed to find the best move at each stage:


1. Generate the whole game tree, down to the leaves.
2. Apply utility (payoff) function to each leaf.
3. Back-up values from leaves through branch nodes:
 - a Max node computes the Max of its child values
 - a Min node computes the Min of its child values
4. At root:
 - If it's a **max node**, choose the move leading to the child with highest value.
 - If it's a **min node**, choose the move leading to the child with lowest value.

Example: a two-ply game tree




- Two symbols (triangles):
 - Max nodes pointing up
 - Min nodes pointing down
- Leaf (terminal) nodes have a utility (payoff) value.

Example: A Two-Ply Game Tree



- The utility of other nodes are determined by applying MinMax algorithm.

Example: A Two-Ply Game Tree



- The best action for Max is A_1
- The best response from Min is A_{11} .

MinMax Algorithm

function **MINIMAX-DECISION**(*state*) returns *an action*

inputs: *state*, current state in game

return the *a* in ACTIONS(*state*) maximizing MIN-VALUE(RESULT(*a*, *state*))

Assuming the root node is Max

function **MAX-VALUE**(*state*) returns *a utility value*

if TERMINAL-TEST(*state*) then return UTILITY(*state*)

$v \leftarrow -\infty$

for *a*, *s* in SUCCESSORS(*state*) do $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function **MIN-VALUE**(*state*) returns *a utility value*

if TERMINAL-TEST(*state*) then return UTILITY(*state*)

$v \leftarrow \infty$

for *a*, *s* in SUCCESSORS(*state*) do $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*

Functions used in the algorithm

Actions(state)

Returns the set of legal moves in the given state state.

Result(action, state)

Returns a state that is the result of applying action action on state state.

Successors(state)

Returns a list of pairs of action-state that can be reached from the given state.

Terminal-Test(state)

Is the game finished? True if finished, false otherwise.

Utility(state)

Gives numerical value of the terminal state state.


Game Tree Size

- Tic-Tac-Toe
 - $b \approx 5$ legal actions per state on average
 - total of 9 moves, $d = 9$
 - $b^d \approx 5^9 = 1,953,125$
 - **Searching the entire tree quite reasonable**
- Chess
 - $b \approx 35$ (approximate average branching factor)
 - $d \approx 100$ (depth of game tree for a “typical” game)
 - $b^d \approx 35^{100} \approx 10^{154}$ nodes!!
 - **Searching the entire tree completely infeasible**

How to reduce the search space

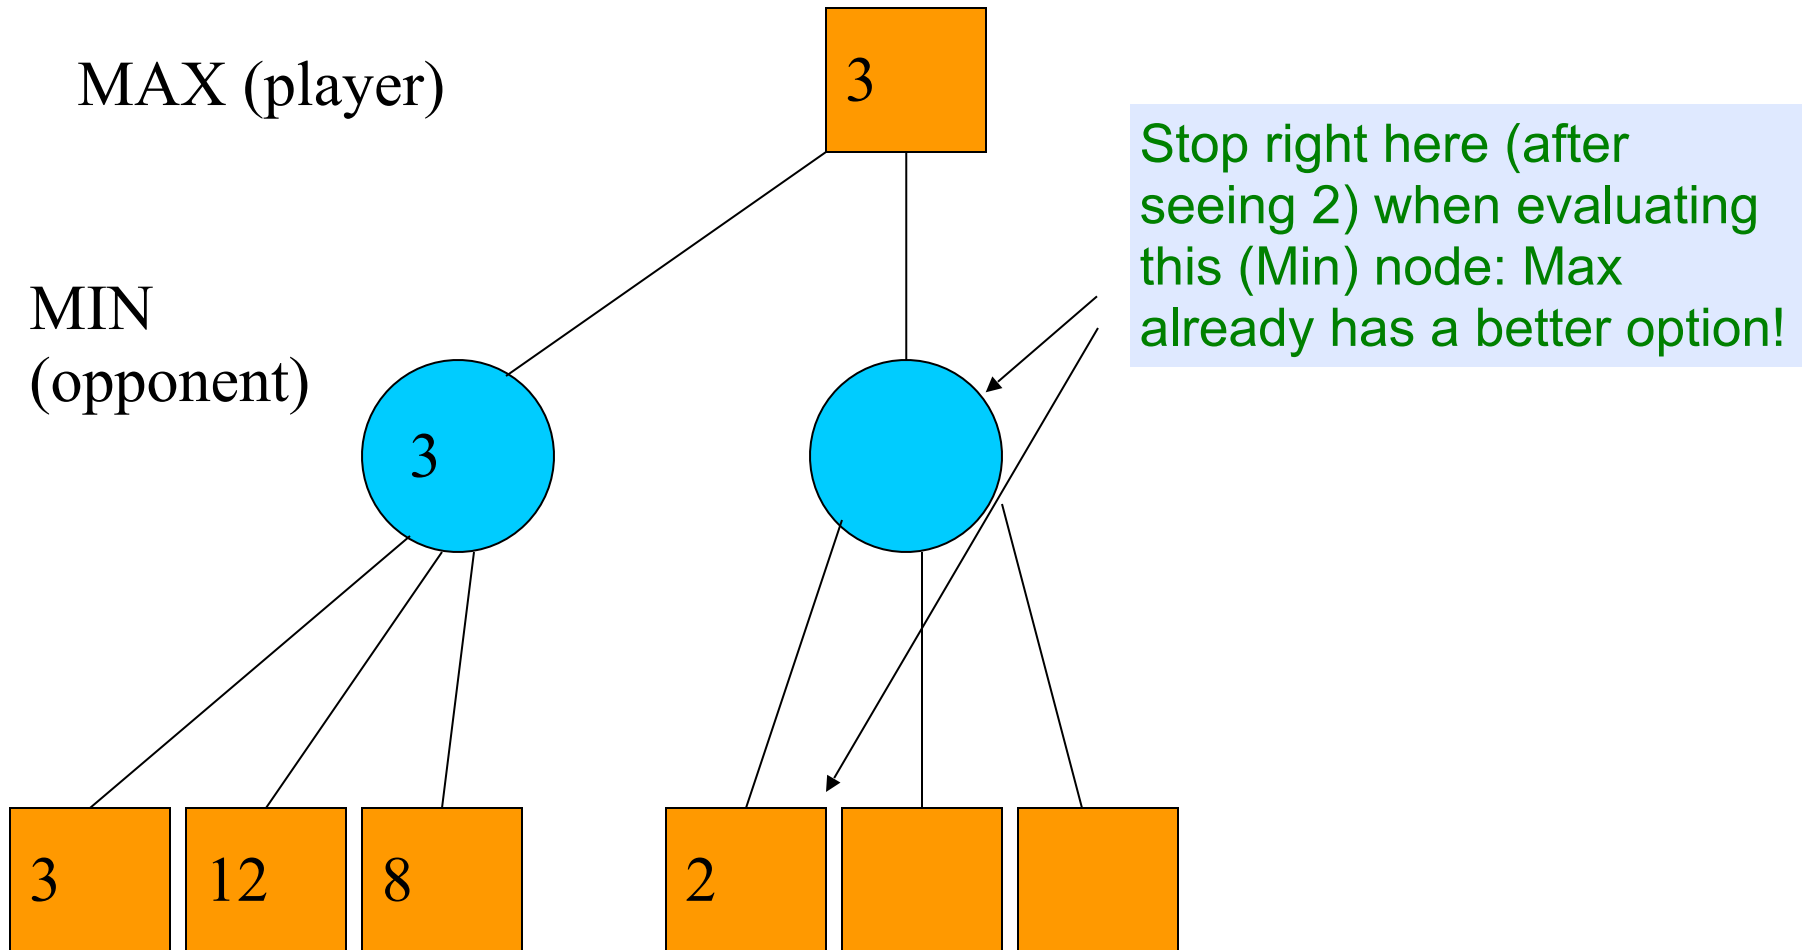
- Pruning the game tree
- Heuristic evaluations of states
- Table lookup instead of search (for opening and closing situations for example)

Pruning: The Idea




If an option is bad compared to other available options, there is no point in expending search time to find out exactly how bad.

Pruning: The Idea




Alpha-Beta Pruning: The Concept



If $m > n$, **Max** would choose the m -node to get a guaranteed utility of at least m .

The Min node with utility n (or less) would never be reached; stop further evaluation.

Alpha-Beta Pruning: The Concept



If $m < n$, **Min** would choose the m -node to get a guaranteed utility of at most m .

The Max node with utility n (or more) would never be reached; stop further evaluation.

Alpha-Beta Pruning Algorithm

Depth first search

- only considers nodes along a single path from root at any time

α = highest-value choice found for MAX higher up in the tree
(initially, $\alpha = -\text{infinity}$)

β = lowest-value choice found for MIN higher up in the tree
(initially, $\beta = +\text{infinity}$)

Pass current values of α and β down to child nodes during search.

Update values of α and β during search:

- MAX updates α at MAX nodes
- MIN updates β at MIN nodes

Prune remaining branches at a node when $\alpha \geq \beta$

Alpha-Beta Pruning Algorithm

function ALPHA-BETA-SEARCH(*state*) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$
 return the *action* in ACTIONS(*state*) with value *v*

Assuming the
root node is Max

function MAX-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
 if $\alpha \geq \beta$ **then return** *v*
 return *v*


function MIN-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow +\infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 $\beta \leftarrow \text{MIN}(\beta, v)$
 if $\alpha \geq \beta$ **then return** *v*
 return *v*

When to Prune


- **Prune whenever $\alpha \geq \beta$**
 - Prune below a Max node whose alpha value becomes greater than or equal to the beta value of its ancestors.
 - **Max nodes update alpha** based on children's returned values.
 - Prune below a Min node whose beta value becomes less than or equal to the alpha value of its ancestors.
 - **Min nodes update beta** based on children's returned values.

Alpha-Beta Example


Do DF-search until first leaf




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)




Alpha-Beta Example (continued)



Alpha-Beta Pruning Example 2



Alpha-Beta Pruning Example 3




Effectiveness of Alpha-Beta Search

- Pruning does not affect the final result (optimal move).
- An entire sub-tree can be pruned.
- Worst-Case
 - branches are ordered so that no pruning takes place. In this case alpha-beta gives no improvement over exhaustive search
- Best-Case
 - each player's best move is the left-most child (i.e., evaluated first)
- Good move *ordering* improves effectiveness of pruning
 - E.g., sort moves by the remembered move values found last time.
 - E.g., expand captures first, then threats, then forward moves, etc.
 - E.g., run Iterative Deepening search, sort by value last iteration.
- In practice often get $O(b^{(d/2)})$ rather than $O(b^d)$
 - this is the same as having a branching factor of \sqrt{b} ,
 - $(\sqrt{b})^d = b^{(d/2)}$, i.e., we effectively go from b to square root of b
 - e.g., in chess go from $b \sim 35$ to $b \sim 6$
 - this permits much deeper search in the same amount of time

Static (Heuristic) Evaluation Functions


- Estimates how good the current board configuration is for a player.
- Typically, evaluate how good it is for the player, how good it is for the opponent, then subtract the opponent's score from the player's.
 - **Othello**: Number of white pieces - Number of black pieces
 - **Chess**: Value of all white pieces - Value of all black pieces
- Typical values from -infinity (loss) to +infinity (win) or $[-1, +1]$.
- If the board evaluation is X for one player, it's $-X$ for the other.
- This allows to perform **cut-off search**: after a maximum depth is reached, use a heuristic evaluation function instead of actual utility.

Evaluation functions



Black to move

White slightly better



White to move

Black winning

For chess, typically *linear* weighted sum of *features*

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g., $w_1 = 9$ with

$f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$